

Credit Card Fraud Detection as a Classification Problem

Business Problem

Credit card companies can detect fraudulent credit card transactions, preventing customers from being charged for products they did not purchase. Data Science and Machine Learning can be used to solve this type of problems. This project aims to demonstrate the modeling of a data set using machine learning with Credit Card Fraud Detection.

```
In [1]: # Importing modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import gridspec
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
In [2]: data = pd.read_csv("creditcard.csv")
data.head().append(data.tail())
```

```
Out[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.36378
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.25542
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.51465
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.38702
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.81773
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.91442
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.58480
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.43245
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.39208
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.48618

10 rows × 31 columns

```
In [3]: display(data.info())
display(data.describe())
display(data.shape)
display(data.isnull().sum())
display(data.duplicated().sum())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1          284807 non-null float64
 2   V2          284807 non-null float64
 3   V3          284807 non-null float64
 4   V4          284807 non-null float64
 5   V5          284807 non-null float64
 6   V6          284807 non-null float64
 7   V7          284807 non-null float64
 8   V8          284807 non-null float64
 9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
None

```

	Time	V1	V2	V3	V4	V5	V6
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552103e-15	2.040130e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01

8 rows × 31 columns

(284807, 31)

```

Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
1081

```

```

In [4]: # drop duplicated values
data.drop_duplicates(inplace =True)

```

Exploration and Visualization

Now we try to find out the relative proportion of valid and fraudulent credit card transactions.

```

In [5]: print("Fraudulent Cases: " + str(len(data[data["Class"] == 1])))
print("Valid Transactions: " + str(len(data[data["Class"] == 0])))
print("Proportion of Fraudulent Cases: " + str(len(data[data["Class"] == 1])/ data.shape[0]))

# To see how small are the number of Fraud transactions
data_pi = data.copy()
data_pi[" "] = np.where(data_pi["Class"] == 1 , "Fraud", "Genuine")

%matplotlib inline
data_pi[" "].value_counts().plot(kind="pie")

```

```


Fraudulent Cases: 473
Valid Transactions: 283253
Proportion of Fraudulent Cases: 0.001667101358352777

```

```

Out[5]: <AxesSubplot:ylabel=' '>

```



Clearly we can see that there is an imbalance in the data with only 0.17% of the total cases are fraudulent.

Check average Money transaction for the fraudulent and no-fraudulent transactions

```
In [6]: print("Average Amount in a Fraudulent Transaction: " + str(data[data["Class"] == 1]["Amount"].mean()))
print("Average Amount in a Valid Transaction: " + str(data[data["Class"] == 0]["Amount"].mean()))
```

Average Amount in a Fraudulent Transaction: 123.87186046511626

Average Amount in a Valid Transaction: 88.41357475466688

As we can clearly notice from this, the average Money transaction for the fraudulent ones are more.


Histograms where values are subgrouped according to Class (valid or fraud)

```
In [7]: # Reorder the columns Amount, Time then the rest
data_plot = data.copy()
amount = data_plot['Amount']
data_plot.drop(labels=['Amount'], axis=1, inplace = True)
data_plot.insert(0, 'Amount', amount)


# Plot the distributions of the features
columns = data_plot.iloc[:,0:30].columns
plt.figure(figsize=(12,30*4))

grids = gridspec.GridSpec(30, 1)
for grid, index in enumerate(data_plot[columns]):
    ax = plt.subplot(grids[grid])
    sns.distplot(data_plot[index][data_plot.Class == 1], hist=False, kde_kws={"shade": True}, bins=100)
    sns.distplot(data_plot[index][data_plot.Class == 0], hist=False, kde_kws={"shade": True}, bins=100)
    ax.set_xlabel("")
    ax.set_title("Distribution of Column: " + str(index))
    ax.legend(labels=['fraudulent', 'valid'])
plt.show()
```


Distribution of Column: Amount




Distribution of Column: Time




Distribution of Column: V1




Distribution of Column: V2














Distribution of Column: V3

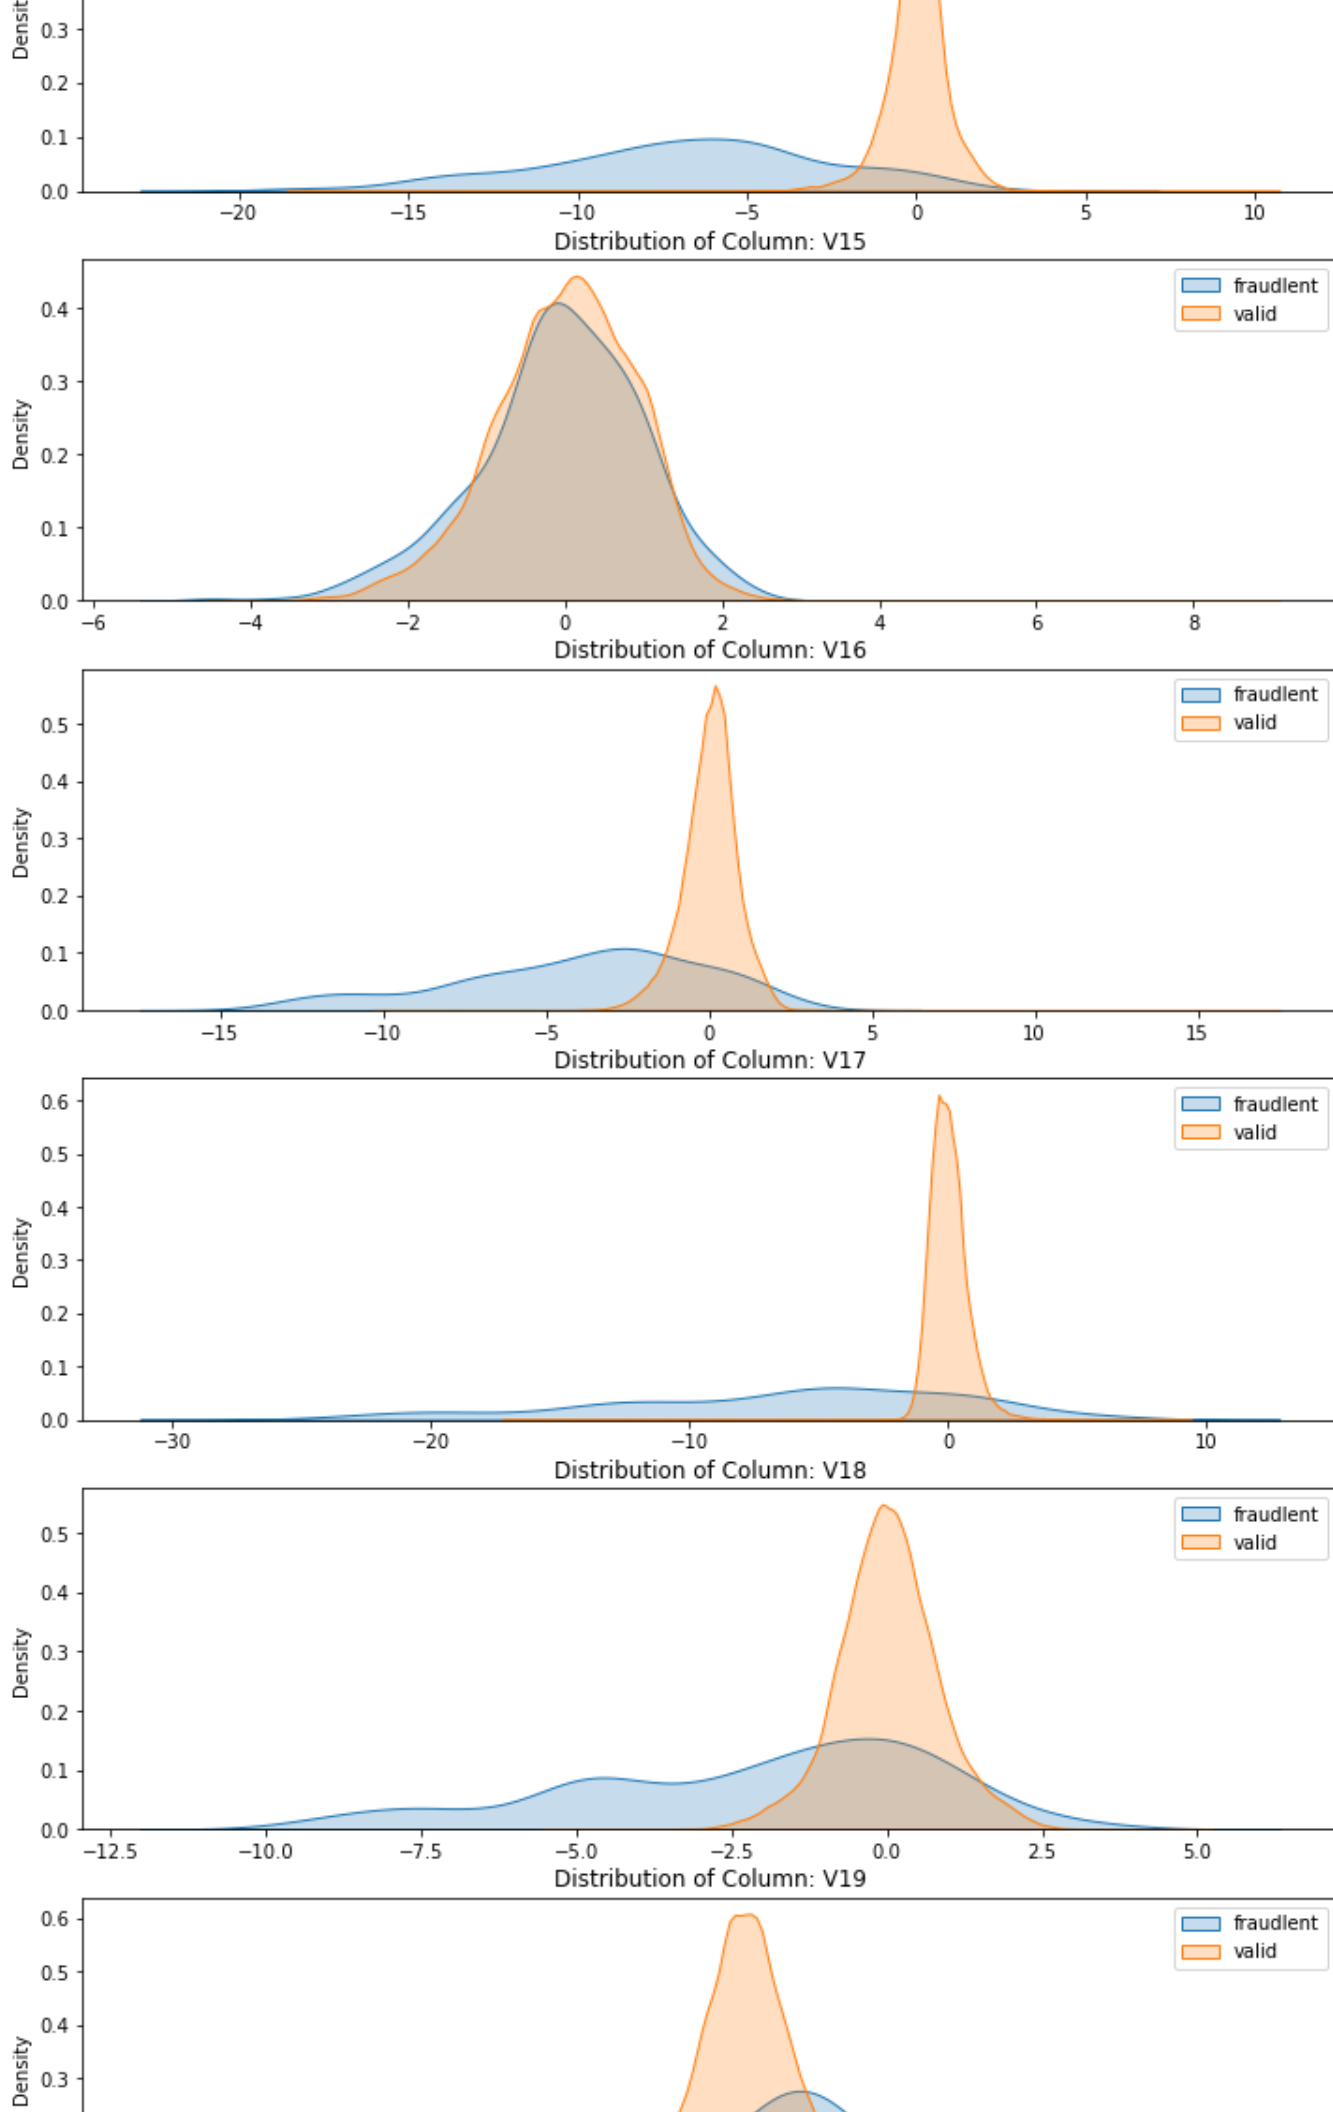


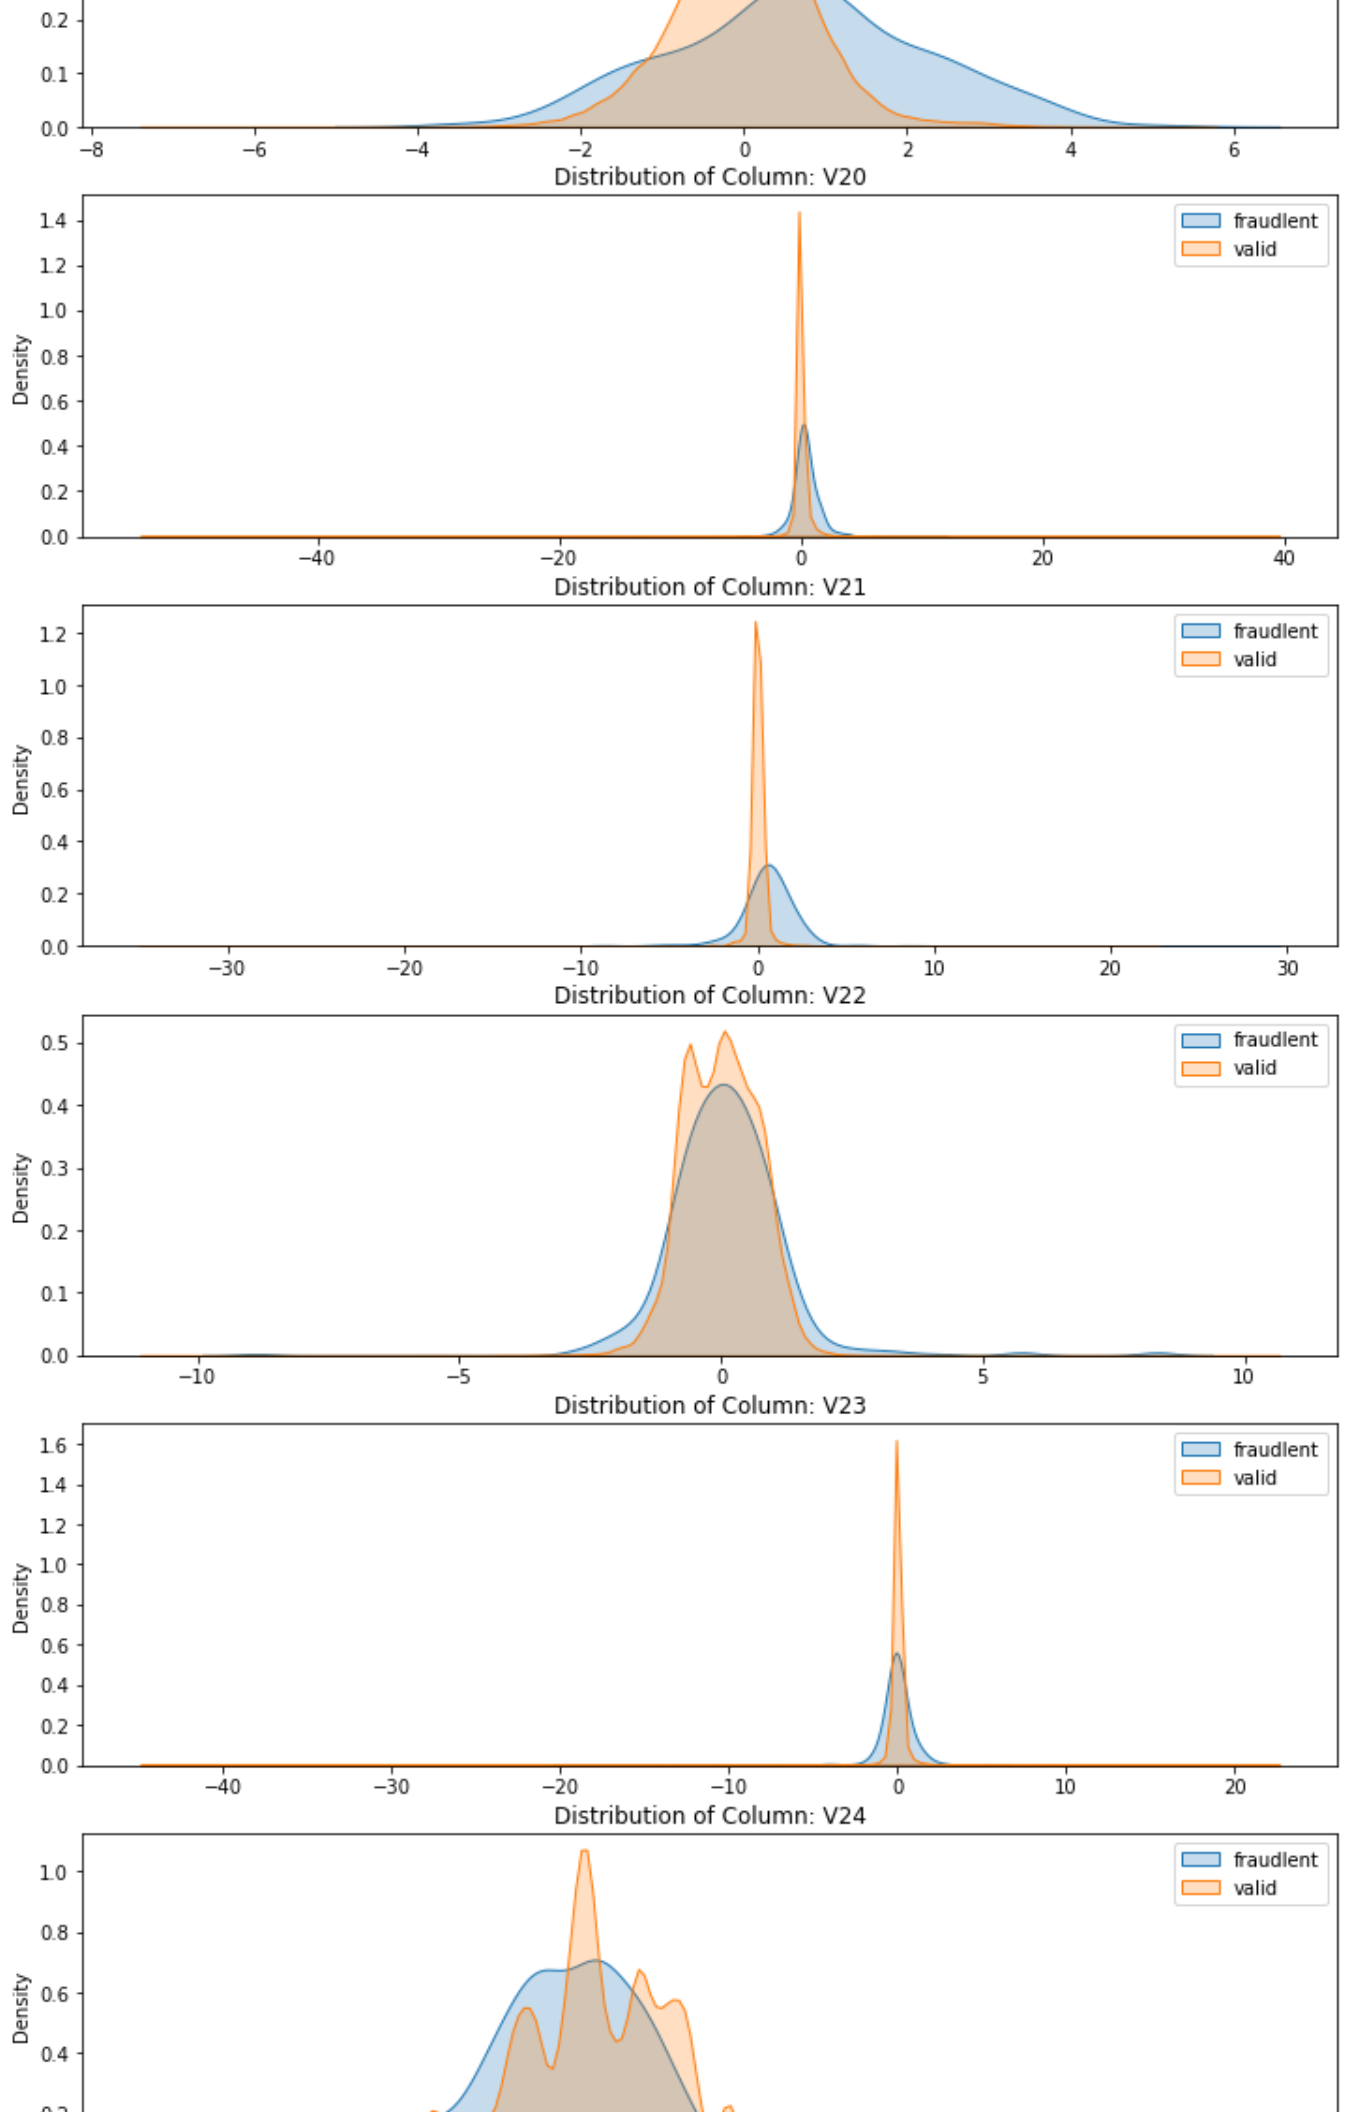
Distribution of Column: V4

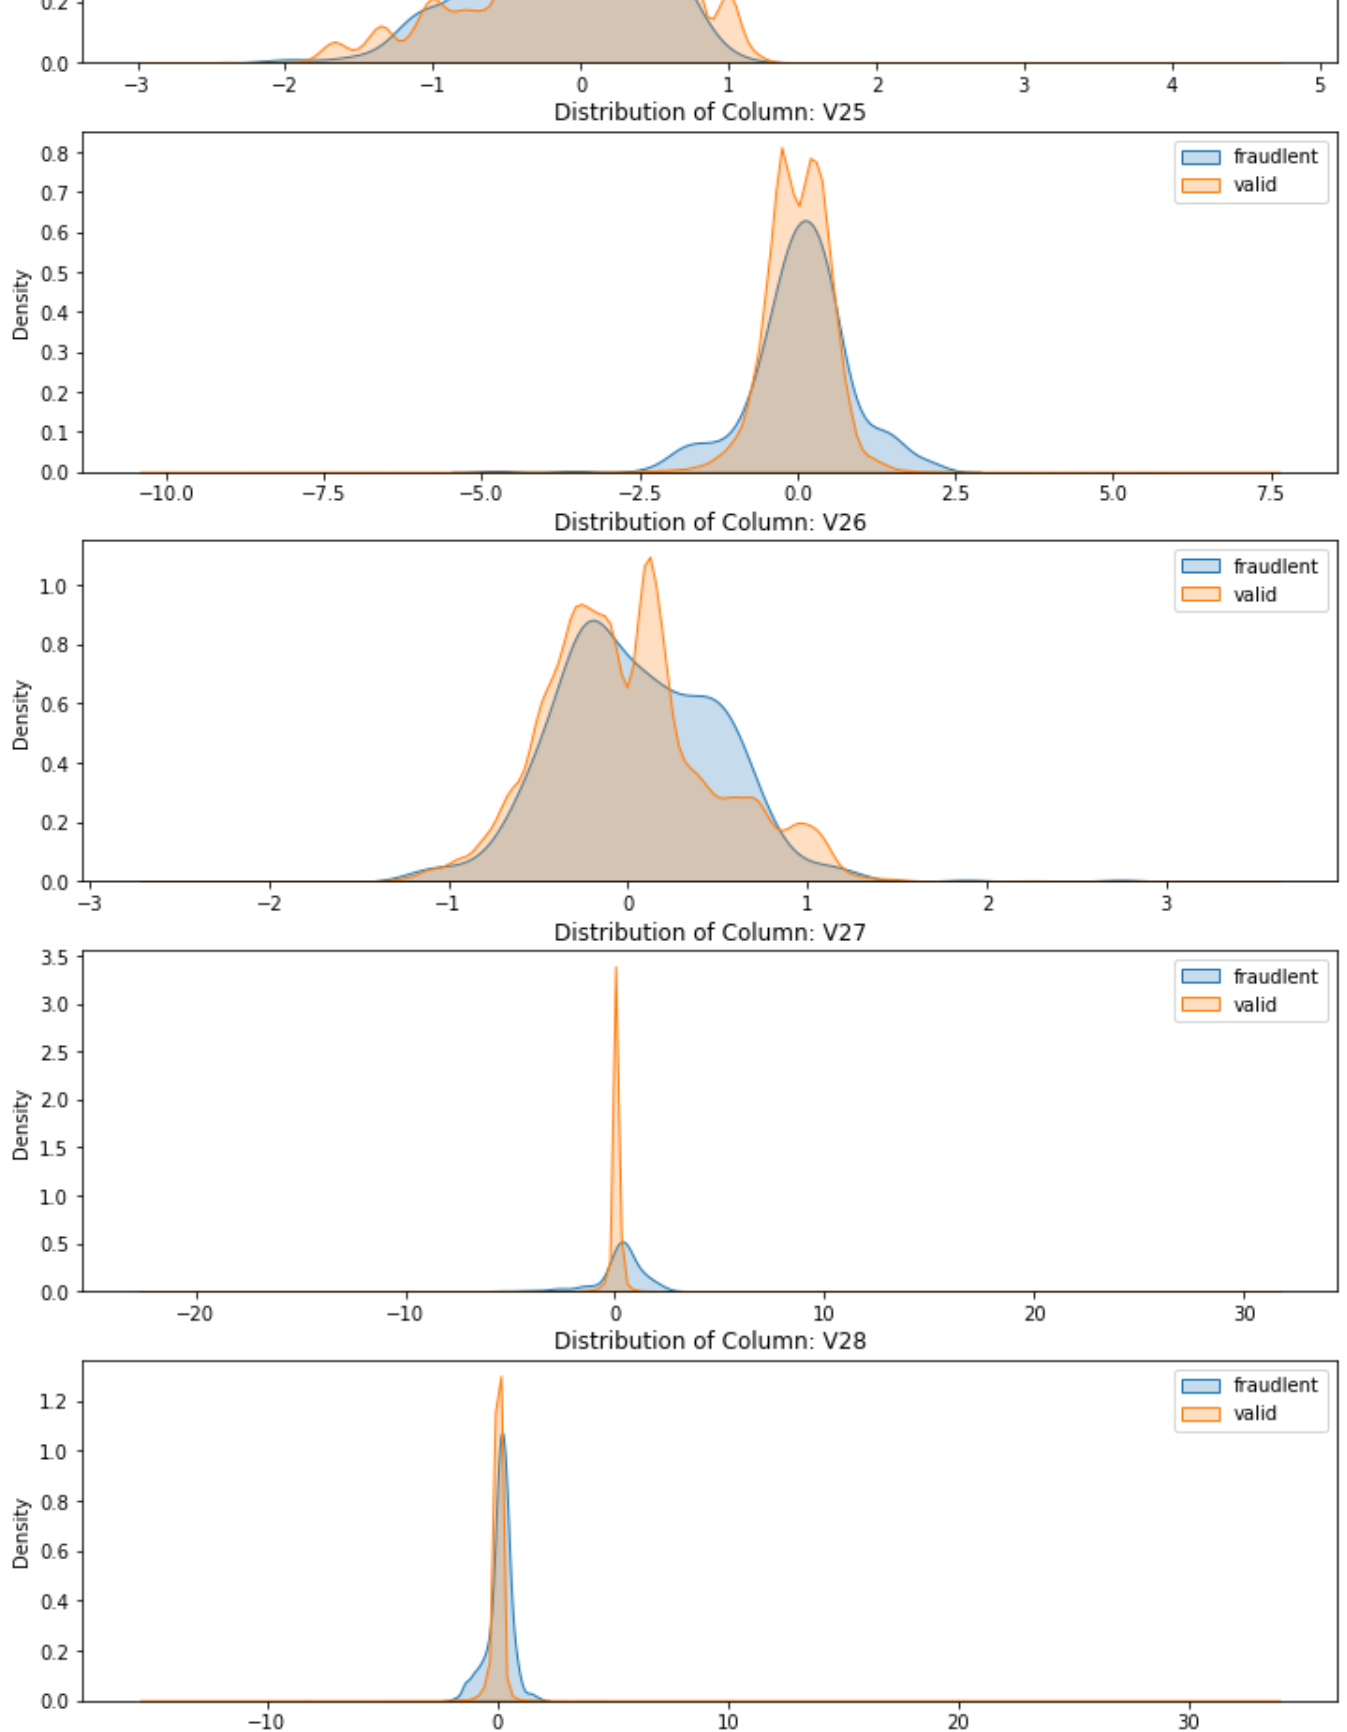












Since there are no missing data, standardization is appropriate. We only use RobustScaler to standardize the Time and Amount columns because all other features of the original dataset from v1 to v28 are obtained using PCA, which is already standardized.

```
In [8]: from sklearn.preprocessing import RobustScaler
scaler = RobustScaler().fit(data[["Time", "Amount"]])
data[["Time", "Amount"]] = scaler.transform(data[["Time", "Amount"]])
```

```
data.head().append(data.tail())
```

Out[8]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	-0.995290	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.3637
1	-0.995290	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.2554
2	-0.995279	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.5146
3	-0.995279	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.3870
4	-0.995267	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.8177
284802	1.035258	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.9144
284803	1.035270	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.5848
284804	1.035282	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.4324
284805	1.035282	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.3920
284806	1.035329	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.4861

10 rows × 31 columns

Modelling

First we divide the data into TARGET and features. And also make the train-test split of the data for further modelling and validation.

In [9]:

```
# Separate TARGET and features
y = data["Class"]
X = data.iloc[:,0:30]

# Use SKLEARN for the split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.2, random_state = 42)
```

Now we describe the flow of the modelling section first and then dive into the sea. As we identified earlier, the dataset is highly imbalanced. Fitting a model on this dataset will result in overfitting towards the majority class. For illustration let's run one model (Random Forest) on the imbalanced data and see the performance.

In [10]:

```
# Using SKLEARN module for random forest
from sklearn.ensemble import RandomForestClassifier

# Fit and predict
naive_rfc = RandomForestClassifier()
naive_rfc.fit(X_train, y_train)
naive_test_preds = naive_rfc.predict(X_test)

# For the performance let's use some metrics from SKLEARN module
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

print("The accuracy is {}".format(accuracy_score(y_test, naive_test_preds)))
print("The precision is {}".format(precision_score(y_test, naive_test_preds)))
```

```
print("The recall is {}".format(recall_score(y_test, naive_test_preds) ))  
print("The f1 score is {}".format(f1_score(y_test, naive_test_preds) ))
```

```
The accuracy is 0.9995418179254926  
The precision is 0.9705882352941176  
The recall is 0.7333333333333333  
The f1 score is 0.8354430379746834
```

One thing to notice here is, we had only 0.17% cases with fraud transactions and a model predicting all transactions to be valid would have similar accuracy. So we need to train our model in a way that is not overfitted to either of the classes. For this, we introduce Oversampling and Undersampling methods. Oversampling resamples from the minority class to balance the class proportions. And undersampling merges or removes similar observations from the majority to achieve the same.

Undersampling

In this section we first describe the structure of the modelling and validations. One trivial point to note is, we will not undersample the test data as we want our model to perform well with skewed class distributions eventually. The steps are as follows (The whole set-up will be structured using the imbalance-learn module):

- Use a 5-fold cross validation on the training set
- On each of the folds use undersampling
- Fit the model on the training folds and validate on the validation fold

```
In [11]: # Create the cross validation framework  
from sklearn.model_selection import StratifiedKFold  
from sklearn.model_selection import GridSearchCV, cross_val_score, RandomizedSearchCV  
  
kf = StratifiedKFold(n_splits=5, random_state = 42, shuffle = True)
```

```
In [12]: #pip install imblearn
```

```
In [13]: # Import the imbalance Learn module  
from imblearn.pipeline import Pipeline, make_pipeline  
from imblearn.under_sampling import NearMiss  
from imblearn.over_sampling import SMOTE  
  
# Import the classifiers  
from sklearn.linear_model import LogisticRegression  
from sklearn.svm import SVC  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier
```

Undersampling - Logistic Regression

```
In [14]: # Logistic Regression  
imba_pipeline = make_pipeline(NearMiss(),  
                              LogisticRegression())  
  
log_reg_params = {"penalty": ['l1', 'l2'],  
                  'C': [ 0.01, 0.1, 1, 100],  
                  'solver' : ['liblinear']}  
  
new_params = {'logisticregression__' + key: log_reg_params[key] for key in log_reg_params}
```

```

grid_imba_log_reg = GridSearchCV(imba_pipeline, param_grid=new_params, cv=kf,
                                return_train_score=True)

grid_imba_log_reg.fit(X_train, y_train);
logistic_cv_score_us = cross_val_score(grid_imba_log_reg, X_train, y_train, scoring = 'recall',

y_test_predict = grid_imba_log_reg.best_estimator_.named_steps['logisticregression'].predict(X_test)
logistic_recall_us = recall_score(y_test, y_test_predict)
logistic_accuracy_us = accuracy_score(y_test, y_test_predict)

log_reg_us = grid_imba_log_reg.best_estimator_

```

```

C:\Users\user\anaconda3\lib\site-packages\sklearn\svm\_base.py:1225: ConvergenceWarning: Liblinear
ar failed to converge, increase the number of iterations.
  warnings.warn(
C:\Users\user\anaconda3\lib\site-packages\sklearn\svm\_base.py:1225: ConvergenceWarning: Liblinear
ar failed to converge, increase the number of iterations.
  warnings.warn(
C:\Users\user\anaconda3\lib\site-packages\sklearn\svm\_base.py:1225: ConvergenceWarning: Liblinear
ar failed to converge, increase the number of iterations.
  warnings.warn(

```

In [15]: log_reg_us, logistic_cv_score_us

```

Out[15]: (Pipeline(steps=[('nearmiss', NearMiss()),
                          ('logisticregression',
                           LogisticRegression(C=0.1, penalty='l1', solver='liblinear'))]),
          array([0.85526316, 0.89473684, 0.92207792, 0.90909091, 0.92207792]))

```

In [16]: log_reg_us, logistic_cv_score_us, logistic_recall_us, logistic_accuracy_us

```

Out[16]: (Pipeline(steps=[('nearmiss', NearMiss()),
                          ('logisticregression',
                           LogisticRegression(C=0.1, penalty='l1', solver='liblinear'))]),
          array([0.85526316, 0.89473684, 0.92207792, 0.90909091, 0.92207792]),
          0.8777777777777778,
          0.742607408451697)

```

In [17]: f1_socre_log = f1_score(y_test, y_test_predict, average = 'weighted')

```

recall_log = recall_score(y_test, y_test_predict)

```

```

precision_log = precision_score(y_test, y_test_predict)

```

```

print(f1_socre_log, recall_log, precision_log)

```

```

0.8507237761700075 0.8777777777777778 0.005383671800463404

```

In [18]: *# Cumulatively create a table for the ROC curve*
from sklearn.metrics import roc_curve, roc_auc_score

```

result_table = pd.DataFrame(columns=['classifiers', 'fpr', 'tpr', 'auc'])

```

```

yproba = grid_imba_log_reg.best_estimator_.named_steps['logisticregression'].predict_proba(X_test)

```

```

fpr, tpr, _ = roc_curve(y_test, yproba)

```

```

auc = roc_auc_score(y_test, yproba)

```

```

result_table = result_table.append({'classifiers': "Logistic Regression",
                                   'fpr':fpr,
                                   'tpr':tpr,

```

```
display(result_table)

'auc':auc}, ignore_index=True)
```

	classifiers	fpr	tpr	auc
0	Logistic Regression	[0.0, 7.060152499293985e-05, 0.000353007624964...	[0.0, 0.0, 0.0, 0.011111111111111112, 0.011111...	0.921291

Undersampling - Random Forest

```
In [19]: # Define the pipeline
imba_pipeline = make_pipeline(NearMiss(),
                               RandomForestClassifier())

params = {
    'n_estimators': [50, 100, 200],
    'max_depth': [4, 6, 10, 12],
    'random_state': [13]
}

new_params = {'randomforestclassifier__' + key: params[key] for key in params}

grid_imba_rf = GridSearchCV(imba_pipeline, param_grid=new_params, cv=kf,
                             return_train_score=True)

grid_imba_rf.fit(X_train, y_train);

rfc_cv_score_us = cross_val_score(grid_imba_rf, X_train, y_train, scoring='recall', cv=kf)

y_test_predict = grid_imba_rf.best_estimator_.named_steps['randomforestclassifier'].predict(X_test)
rfc_recall_us = recall_score(y_test, y_test_predict)
rfc_accuracy_us = accuracy_score(y_test, y_test_predict)

rfc = grid_imba_rf.best_estimator_
```

```
In [20]: rfc, rfc_recall_us, rfc_accuracy_us, rfc_cv_score_us
```

```
Out[20]: (Pipeline(steps=[('nearmiss', NearMiss()),
                           ('randomforestclassifier',
                            RandomForestClassifier(max_depth=4, n_estimators=50,
                                                    random_state=13))]),
          0.9555555555555556,
          0.189669756458605,
          array([0.93421053, 0.96052632, 0.94805195, 0.96103896, 1.          ]))
```

```
In [21]: f1_socre_rfc = f1_score(y_test, y_test_predict, average = 'weighted')

recall_rfc= recall_score(y_test, y_test_predict)

precision_rfc = precision_score(y_test, y_test_predict)

print(f1_socre_rfc, recall_rfc, precision_rfc)

0.3166242961086456 0.9555555555555556 0.0018669271681319875
```

```
In [22]: # Cumulatively create a table for the ROC curve
yproba = grid_imba_rf.best_estimator_.named_steps['randomforestclassifier'].predict_proba(X_test)

fpr, tpr, _ = roc_curve(y_test, yproba)
auc = roc_auc_score(y_test, yproba)
```

```
result_table = result_table.append({'classifiers': "Random Forest",
                                   'fpr':fpr,
                                   'tpr':tpr,
                                   'auc':auc}, ignore_index=True)

display(result_table)
```

	classifiers	fpr	tpr	auc
0	Logistic Regression	[0.0, 7.060152499293985e-05, 0.000353007624964...	[0.0, 0.0, 0.0, 0.011111111111111112, 0.011111...	0.921291
1	Random Forest	[0.0, 8.825190624117481e-05, 8.825190624117481...	[0.0, 0.4444444444444444, 0.5222222222222223, ...	0.873910

Undersampling - Support Vector Classifier

```
In [23]: # Define the pipeline
imba_pipeline = make_pipeline(NearMiss(),
                              SVC(probability = True))
svc_params = {'C': [0.5, 0.7, 0.9, 1],
              'kernel': ['rbf', 'poly', 'sigmoid', 'linear']}

new_params = {'svc__' + key: svc_params[key] for key in svc_params}

grid_imba_svc = GridSearchCV(imba_pipeline, param_grid=new_params, cv=kf,
                             return_train_score=True)

grid_imba_svc.fit(X_train, y_train);

svc_cv_score_us = cross_val_score(grid_imba_svc, X_train, y_train, scoring='recall', cv=kf)

y_test_predict = grid_imba_svc.best_estimator_.named_steps['svc'].predict(X_test)
svc_recall_us = recall_score(y_test, y_test_predict)
svc_accuracy_us = accuracy_score(y_test, y_test_predict)

svc = grid_imba_svc.best_estimator_
```

```
In [24]: svc, svc_recall_us, svc_accuracy_us, svc_cv_score_us
```

```
Out[24]: (Pipeline(steps=[('nearmiss', NearMiss()),
                          ('svc', SVC(C=0.5, kernel='poly', probability=True))]),
 0.6,
 0.9916117435590174,
 array([0.63157895, 0.59210526, 0.75324675, 0.71428571, 0.62337662]))
```

```
In [25]: f1_socre_svc = f1_score(y_test, y_test_predict, average = 'weighted')

recall_svc = recall_score(y_test, y_test_predict)

precision_svc = precision_score(y_test, y_test_predict)

print(f1_socre_svc, recall_svc, precision_svc)

0.9944981538721104 0.6 0.10931174089068826
```

```
In [26]: # Cumulatively create a table for the ROC curve
yproba = grid_imba_svc.best_estimator_.named_steps['svc'].predict_proba(X_test)[::,1]
```

```
fpr, tpr, _ = roc_curve(y_test, yproba)
auc = roc_auc_score(y_test, yproba)

result_table = result_table.append({'classifiers': "Support Vector Classifier",
                                     'fpr':fpr,
                                     'tpr':tpr,
                                     'auc':auc}, ignore_index=True)

display(result_table)
```

	classifiers		fpr	tpr	auc
0	Logistic Regression	[0.0, 7.060152499293985e-05, 0.000353007624964...	[0.0, 0.0, 0.0, 0.011111111111111112, 0.011111...	0.921291	
1	Random Forest	[0.0, 8.825190624117481e-05, 8.825190624117481...	[0.0, 0.4444444444444444, 0.5222222222222223, ...	0.873910	
2	Support Vector Classifier	[0.0, 0.0011825755436317424, 0.001217876306128...	[0.0, 0.2222222222222222, 0.222222222222222, ...	0.958269	

Undersampling - Decision Tree Classifier

```
In [27]: # DecisionTree Classifier
imba_pipeline = make_pipeline(NearMiss(),
                              DecisionTreeClassifier())

tree_params = {"criterion": ["gini", "entropy"], "max_depth": list(range(2,4,1)),
               "min_samples_leaf": list(range(5,7,1))}
new_params = {'decisiontreeclassifier__' + key: tree_params[key] for key in tree_params}
#grid_imba_tree = GridSearchCV(imba_pipeline, param_grid=new_params, cv=kf, scoring='recall',
#                               return_train_score=True)
grid_imba_tree = GridSearchCV(imba_pipeline, param_grid=new_params, cv=kf,
                              return_train_score=True)

grid_imba_tree.fit(X_train, y_train);
dtree_cv_score_us = cross_val_score(grid_imba_tree, X_train, y_train, scoring='recall', cv=kf)

y_test_predict = grid_imba_tree.best_estimator_.named_steps['decisiontreeclassifier'].predict(X_test)
dtree_recall_us = recall_score(y_test, y_test_predict)
dtree_accuracy_us = accuracy_score(y_test, y_test_predict)

# print("Cross Validation Score for Decision Tree Classifier: " + str(udtree_cv_score.mean()))
# print("Recall Score for Decision Tree Classifier: " + str(udtree_recall))
tree_clf = grid_imba_tree.best_estimator_
```

```
In [28]: tree_clf, dtree_accuracy_us, dtree_recall_us, dtree_cv_score_us
```

```
Out[28]: (Pipeline(steps=[('nearmiss', NearMiss()),
                          ('decisiontreeclassifier',
                           DecisionTreeClassifier(max_depth=2, min_samples_leaf=5))]),
 0.6681704437317167,
 0.8222222222222222,
 array([0.89473684, 0.90789474, 0.94805195, 0.8961039 , 0.83116883]))
```

```
In [29]: f1_socre_dtree = f1_score(y_test, y_test_predict, average = 'weighted')

recall_dtree = recall_score(y_test, y_test_predict)

precision_dtree = precision_score(y_test, y_test_predict)
```



```
print(f1_socre_dtree, recall_dtree, precision_dtree)

0.7995125912644028 0.8222222222222222 0.003917831427361288
```

```
In [30]: # Cumulatively create a table for the ROC curve
yproba = grid_imba_tree.best_estimator_.named_steps['decisiontreeclassifier'].predict_proba(X_test)

fpr, tpr, _ = roc_curve(y_test, yproba)
auc = roc_auc_score(y_test, yproba)

result_table = result_table.append({'classifiers': "Decision Tree",
                                     'fpr':fpr,
                                     'tpr':tpr,
                                     'auc':auc}, ignore_index=True)

display(result_table)
```

	classifiers		fpr	tpr	auc
0	Logistic Regression	[0.0, 7.060152499293985e-05, 0.000353007624964...	[0.0, 0.0, 0.0, 0.011111111111111112, 0.011111...	0.921291	
1	Random Forest	[0.0, 8.825190624117481e-05, 8.825190624117481...	[0.0, 0.44444444444444444, 0.5222222222222223, ...	0.873910	
2	Support Vector Classifier	[0.0, 0.0011825755436317424, 0.001217876306128...	[0.0, 0.2222222222222222, 0.222222222222222, ...	0.958269	
3	Decision Tree	[0.0, 0.24634637108161536, 0.3243610561988139,...	[0.0, 0.02222222222222223, 0.822222222222222...	0.650573	

Undersampling - k-Nearest Neighbour Classifier

```
In [31]: # KNeighbors Classifier
imba_pipeline = make_pipeline(NearMiss(),
                               KNeighborsClassifier())

knears_params = {"n_neighbors": list(range(2,5,1)),
                 'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute']}

new_params = {'kneighborsclassifier__' + key: knears_params[key] for key in knears_params}

grid_imba_knn = GridSearchCV(imba_pipeline, param_grid=new_params, cv=kf,
                             return_train_score=True)

grid_imba_knn.fit(X_train, y_train);
knear_cv_score_us = cross_val_score(grid_imba_knn, X_train, y_train, scoring='recall', cv=kf)

y_test_predict = grid_imba_knn.best_estimator_.named_steps['kneighborsclassifier'].predict(X_test)
knear_recall_us = recall_score(y_test, y_test_predict)
knear_accuracy_us = accuracy_score(y_test, y_test_predict)

knears_neighbors = grid_imba_knn.best_estimator_
```

```
In [32]: knears_neighbors, knear_accuracy_us, knear_recall_us, knear_cv_score_us
```

```
Out[32]: (Pipeline(steps=[('nearmiss', NearMiss()),
                           ('kneighborsclassifier', KNeighborsClassifier(n_neighbors=4))]),
          0.8127797554012618,
          0.8666666666666667,
          array([0.84210526, 0.93421053, 0.90909091, 0.93506494, 0.8961039 ]))
```

```
In [33]: f1_socre_knears = f1_score(y_test, y_test_predict, average = 'weighted')
```

```
recall_knears= recall_score(y_test, y_test_predict)
```

```
precision_knears= precision_score(y_test, y_test_predict)
```

```
print(f1_socre_knears, recall_knears, precision_knears)
```

```
0.8951661389883894 0.8666666666666667 0.007296538821328344
```

```
In [34]: # Cumulatively create a table for the ROC curve
```

```
yproba = grid_imba_knn.best_estimator_.named_steps['kneighborsclassifier'].predict_proba(X_test)
```

```
fpr, tpr, _ = roc_curve(y_test, yproba)
```

```
auc = roc_auc_score(y_test, yproba)
```

```
result_table = result_table.append({'classifiers': "k-Nearest Neighbour",
                                     'fpr':fpr,
                                     'tpr':tpr,
                                     'auc':auc}, ignore_index=True)
```

```
display(result_table)
```

	classifiers		fpr	tpr	auc
0	Logistic Regression	[0.0, 7.060152499293985e-05, 0.000353007624964...	[0.0, 0.0, 0.0, 0.011111111111111112, 0.011111...	0.921291	
1	Random Forest	[0.0, 8.825190624117481e-05, 8.825190624117481...	[0.0, 0.44444444444444444, 0.5222222222222223, ...	0.873910	
2	Support Vector Classifier	[0.0, 0.0011825755436317424, 0.001217876306128...	[0.0, 0.2222222222222222, 0.222222222222222, ...	0.958269	
3	Decision Tree	[0.0, 0.24634637108161536, 0.3243610561988139,...	[0.0, 0.02222222222222223, 0.822222222222222...	0.650573	
4	k-Nearest Neighbour	[0.0, 0.07884425303586558, 0.1873058458062694,...	[0.0, 0.8444444444444444, 0.8666666666666667, ...	0.885104	

Summarize the undersampling model performances

```
In [35]: # Gather the scores
```

```
data_score = [['Logistic Regression', logistic_cv_score_us.mean(), logistic_accuracy_us, logistic_recall_us],
               ['Random Forest', rfc_cv_score_us.mean(), rfc_accuracy_us, rfc_recall_us],
               ['Support Vector', svc_cv_score_us.mean(), svc_accuracy_us, svc_recall_us],
               ['Decision Tree', dtree_cv_score_us.mean(), dtree_accuracy_us, dtree_recall_us],
               ['k-Nearest Neighbour', knear_cv_score_us.mean(), knear_accuracy_us, knear_recall_us]]
```

```
# Create the dataframe
```

```
data_table = pd.DataFrame(data_score, columns = ['Classifier', 'CV Score', 'Accuracy', 'Recall Score'])
data_table
```

Out[35]:

	Classifier	CV Score	Accuracy	Recall Score
0	Logistic Regression	0.900649	0.742607	0.877778
1	Random Forest	0.960766	0.189670	0.955556
2	Support Vector	0.662919	0.991612	0.600000
3	Decision Tree	0.895591	0.668170	0.822222
4	k-Nearest Neighbour	0.903315	0.812780	0.866667

```
In [36]: # Plot the ROC curve for undersampling
result_table.set_index('classifiers', inplace=True)
fig = plt.figure(figsize=(16,6))

for i in result_table.index:
    plt.plot(result_table.loc[i]['fpr'],
             result_table.loc[i]['tpr'],
             label="{}, AUC={:.3f}".format(i, result_table.loc[i]['auc']))

plt.plot([0,1], [0,1], color='orange', linestyle='--')

plt.xticks(np.arange(0.0, 1.1, step=0.1))
plt.xlabel("Flase Positive Rate", fontsize=15)

plt.yticks(np.arange(0.0, 1.1, step=0.1))
plt.ylabel("True Positive Rate", fontsize=15)

plt.title('ROC Curve Analysis for Undersampling', fontweight='bold', fontsize=15)
plt.legend(prop={'size':13}, loc='lower right')

plt.show()
```

